

# Package ‘dbscan’

December 15, 2015

**Version** 0.9-6

**Date** 2015-12-14

**Title** Density Based Clustering of Applications with Noise (DBSCAN) and Related Algorithms

**Description** A fast reimplementa-tion of several density-based algorithms of the DBSCAN family for spatial data. Includes the DBSCAN (density-based spatial clustering of applications with noise) and OPTICS (ordering points to identify the clustering structure) clustering algorithms and the LOF (local outlier factor) algorithm. The implementation uses the kd-tree data structure (from library ANN) for faster k-nearest neighbor search. An R interface to fast kNN and fixed-radius NN search is also provided.

**Imports** Rcpp, graphics, stats, methods

**LinkingTo** Rcpp

**Suggests** fpc, microbenchmark, testthat

**BugReports** <https://github.com/mhahsler/dbscan/issues>

**License** GPL (>= 2)

**Copyright** ANN library is copyright by University of Maryland, Sunil Arya and David Mount. All other code is copyright by Michael Hahsler.

**NeedsCompilation** yes

**Author** Michael Hahsler [aut, cre, cph],  
Sunil Arya [ctb, cph],  
David Mount [ctb, cph]

**Maintainer** Michael Hahsler <mhahsler@lyle.smu.edu>

**Repository** CRAN

**Date/Publication** 2015-12-15 01:14:05

## R topics documented:

dbscan . . . . .	2
frNN . . . . .	4
kNN . . . . .	6

kNNdist . . . . .	8
lof . . . . .	9
optics . . . . .	10

<b>Index</b>	<b>13</b>
--------------	-----------

---

dbscan	<i>DBSCAN</i>
--------	---------------

---

**Description**

Fast reimplementa-tion of the DBSCAN (Density-based spatial clustering of applications with noise) clustering algorithm using a kd-tree. The implementation is significantly faster and can work with larger data sets than dbscan in **fpc**.

**Usage**

```
dbscan(x, eps, minPts = 5, weights = NULL,  
       borderPoints = TRUE, search = "kdtree", bucketSize = 10,  
       splitRule = "suggest", approx = 0)
```

**Arguments**

x	a data matrix or a dist object.
eps	size of the epsilon neighborhood.
minPts	number of minimum points in the eps region (for core points). Default is 5 points.
weights	numeric; weights for the data points. Only needed to perform weighted clustering.
borderPoints	logical; should border points be assigned. The default is TRUE for regular DBSCAN. If FALSE then border points are considered noise (see DBSCAN* in Campello et al, 2013).
search	nearest neighbor search strategy (one of "kdtree" or "linear", "dist").
bucketSize	max size of the kd-tree leafs.
splitRule	rule to split the kd-tree. One of "STD", "MIDPT", "FAIR", "SL_MIDPT", "SL_FAIR" or "SUGGEST" (SL stands for sliding). "SUGGEST" uses ANNs best guess.
approx	relative error bound for approximate nearest neighbor searching.

**Details**

This implementation of DBSCAN implements the original algorithm as described by Ester et al (1996). DBSCAN estimates the density around each data point by counting the number of points in a user-specified eps-neighborhood and applies a user-specified minPts thresholds to identify core, border and noise points. In a second step, core points are joined into a cluster if they are density-reachable (i.e., there is a chain of core points where one falls inside the eps-neighborhood of the

next). Finally, border points are assigned to clusters. The algorithm only needs parameters `eps` and `minPts`.

Border points are arbitrarily assigned to clusters in the original algorithm. DBSCAN\* (see Campello et al 2013) treats all border points as noise points. This is implemented with `borderPoints = FALSE`.

Setting parameters for DBSCAN: `minPts` is often set to be dimensionality of the data plus one. The knee in [kNNdistplot](#) can be used to find suitable values for `eps`.

See [kNN](#) for more information on the other parameters related to nearest neighbor search.

### Value

A object of class 'dbscan' with the following components:

<code>eps</code>	value of the <code>eps</code> parameter.
<code>minPts</code>	value of the <code>minPts</code> parameter.
<code>cluster</code>	A integer vector with cluster assignments. Zero indicates noise points.

### Author(s)

Michael Hahsler

### References

Martin Ester, Hans-Peter Kriegel, Joerg Sander, Xiaowei Xu (1996). A Density-Based Algorithm for Discovering Clusters in Large Spatial Databases with Noise. Institute for Computer Science, University of Munich. *Proceedings of 2nd International Conference on Knowledge Discovery and Data Mining (KDD-96)*.

Campello, R. J. G. B.; Moulavi, D.; Sander, J. (2013). Density-Based Clustering Based on Hierarchical Density Estimates. *Proceedings of the 17th Pacific-Asia Conference on Knowledge Discovery in Databases, PAKDD 2013*, Lecture Notes in Computer Science 7819, p. 160.

### See Also

[kNNdistplot](#), [dbscan](#) in [fpc](#).

### Examples

```
data(iris)
iris <- as.matrix(iris[,1:4])

res <- dbscan(iris, eps = .4, minPts = 4)
res

pairs(iris, col = res$cluster + 1L)

## compare with dbscan from package fpc (only if installed)
if (requireNamespace("fpc", quietly = TRUE)) {
  res2 <- fpc::dbscan(iris, eps = .4, MinPts = 4)
  pairs(iris, col = res2$cluster + 1L)
```

```

## make sure both version produce the same results
all(res$cluster == res2$cluster)
}

## find suitable eps parameter (look at knee)
kNNdistplot(iris, k = 4)

## example data from fpc
set.seed(665544)
n <- 100
x <- cbind(
  x = runif(10, 0, 10) + rnorm(n, sd = 0.2),
  y = runif(10, 0, 10) + rnorm(n, sd = 0.2)
)

res <- dbscan::dbscan(x, eps = .2, minPts = 4)
res

plot(x, col=res$cluster + 1L)

## Not run:
## compare speed against fpc version (if microbenchmark is installed)
if (requireNamespace("microbenchmark", quietly = TRUE)) {
  t_dbscan <- microbenchmark::microbenchmark(
    dbscan::dbscan(x, .2, 4), times = 10, unit = "ms")
  t_dbscan_linear <- microbenchmark::microbenchmark(
    dbscan::dbscan(x, .2, 4, search = "linear"), times = 10, unit = "ms")
  t_fpc <- microbenchmark::microbenchmark(
    fpc::dbscan(x, .2, 4), times = 10, unit = "ms")

  rbind(t_fpc, t_dbscan_linear, t_dbscan)

  boxplot(rbind(t_fpc, t_dbscan_linear, t_dbscan),
    names = c("fpc (R)", "dbscan (linear)", "dbscan (kd-tree)"),
    main = "Runtime comparison in ms")

  ## speedup of the kd-tree-based version compared to the fpc implementation
  median(t_fpc$time) / median(t_dbscan$time)
}

## End(Not run)

```

## Description

This function uses a kd-tree to find the fixed radius nearest neighbors (including distances) fast.

**Usage**

```
frNN(x, eps, sort = TRUE, search = "kdtree", bucketSize = 10,  
     splitRule = "suggest", approx = 0)
```

**Arguments**

x	a data matrix or a dist object.
eps	neighbors radius.
search	nearest neighbor search strategy (one of "kdtree" or "linear", "dist").
sort	sort the neighbors by distance?
bucketSize	max size of the kd-tree leafs.
splitRule	rule to split the kd-tree. One of "STD", "MIDPT", "FAIR", "SL_MIDPT", "SL_FAIR" or "SUGGEST" (SL stands for sliding). "SUGGEST" uses ANNs best guess.
approx	use approximate nearest neighbors. All NN up to a distance of a factor of 1+approx eps may be used. Some actual NN may be omitted leading to spurious clusters and noise points. However, the algorithm will enjoy a significant speedup.

**Details**

For details on the parameters see [kNN](#).

Note: self-matches are not returned!

**Value**

A list with the following components:

dist	a matrix with distances.
id	a matrix with ids.
eps	eps used.

**Author(s)**

Michael Hahsler

**References**

David M. Mount and Sunil Arya (2010). ANN: A Library for Approximate Nearest Neighbor Searching, <https://www.cs.umd.edu/~mount/ANN/>.

**See Also**

[kNN](#) for k nearest neighbor search.

## Examples

```
data(iris)

# Find fixed radius nearest neighbors for each point
nn <- frNN(iris[,-5], eps=.5)

# Number of neighbors
hist(sapply(nn$id, length),
     xlab = "k", main="Number of Neighbors",
     sub = paste("Neighborhood size eps =", nn$eps))

# Explore neighbors of point i = 10
i <- 10
nn$id[[i]]
nn$dist[[i]]

plot(iris[,-5], col = ifelse(1:nrow(iris) %in% nn$id[[i]], "red", "black"))
```

---

kNN

*Find the k Nearest Neighbors*


---

## Description

This function uses a kd-tree to find all k nearest neighbors in a data matrix (including distances) fast.

## Usage

```
kNN(x, k, sort = TRUE, search = "kdtree", bucketSize = 10,
    splitRule = "suggest", approx = 0)
```

## Arguments

x	a data matrix or a dist object.
k	number of neighbors to find.
search	nearest neighbor search strategy (one of "kdtree", "linear" or "dist").
sort	sort the neighbors by distance?
bucketSize	max size of the kd-tree leafs.
splitRule	rule to split the kd-tree. One of "STD", "MIDPT", "FAIR", "SL_MIDPT", "SL_FAIR" or "SUGGEST" (SL stands for sliding). "SUGGEST" uses ANNs best guess.
approx	use approximate nearest neighbors. All NN up to a distance of a factor of 1+approx eps may be used. Some actual NN may be omitted leading to spurious clusters and noise points. However, the algorithm will enjoy a significant speedup.

## Details

If `x` is a matrix then this implementation uses the ANN library (see Mount and Arya, 2010) for nearest neighbor search. Linear nearest neighbor search can be used. To speed up nearest neighbor search the kd-tree is used.

Note: self-matches are removed!

`bucketSize` and `splitRule` influence how the kd-tree is built. `approx` uses the approximate nearest neighbor search implemented in ANN. All nearest neighbors up to a distance of  $\text{eps}/(1+\text{approx})$  will be considered and all with a distance greater than `eps` will not be considered. The other points might be considered. Note that this results in some actual nearest neighbors being omitted leading to spurious clusters and noise points. However, the algorithm will enjoy a significant speedup. For more details see Mount and Arya (2010).

## Value

A list with the following components:

<code>dist</code>	a matrix with distances.
<code>id</code>	a matrix with ids.
<code>k</code>	number of <code>k</code> used.

## Author(s)

Michael Hahsler

## References

David M. Mount and Sunil Arya (2010). ANN: A Library for Approximate Nearest Neighbor Searching, <http://www.cs.umd.edu/~mount/ANN/>.

## See Also

[frNN](#) for fixed radius nearest neighbors.

## Examples

```
data(iris)
# finding kNN directly in data (using a kd-tree)
nn <- kNN(iris[,-5], k=10)
nn

# explore neighborhood of point 10
i <- 10
nn$id[i,]
plot(iris[,-5], col = ifelse(1:nrow(iris) %in% nn$id[i,], "red", "black"))
```

---

kNNdist*Calculate and plot the k-Nearest Neighbor Distance*

---

**Description**

Fast calculation of the k-nearest neighbor distances in a matrix of points. The plot can be used to help find a suitable value for the eps neighborhood for DBSCAN. Look for the knee in the plot.

**Usage**

```
kNNdist(x, k, ...)  
kNNdistplot(x, k = 4, ...)
```

**Arguments**

x	the data set as a matrix or a dist object.
k	number of nearest neighbors used (use minPoints).
...	further arguments are passed on to kNN.

**Details**

See [kNN](#) for a discussion of the kd-tree related parameters.

**Value**

kNNdist returns a numeric vector with the distance to its k nearest neighbor.

**Author(s)**

Michael Hahsler

**See Also**

[kNN](#).

**Examples**

```
data(iris)  
iris <- as.matrix(iris[,1:4])  
  
kNNdist(iris, k=4, search="kd")  
kNNdistplot(iris, k=4)  
## the knee is around a distance of .5  
  
cl <- dbscan(iris, eps = .5, minPts = 4)  
pairs(iris, col = cl$cluster+1L)  
## Note: black are noise points
```



---

lof	<i>Local Outlier Factor Score</i>
-----	-----------------------------------

---

**Description**

Calculate the Local Outlier Factor (LOF) score for each data point using a kd-tree to speed up kNN search.

**Usage**

```
lof(x, k = 4, ...)
```

**Arguments**

x	a data matrix or a dist object.
k	size if neighborhood.
...	further arguments are passed on to kNN.

**Details**

LOF compares the local density of an point to the local densities of its neighbors. Points that have a substantially lower density than their neighbors are considered outliers. A LOF score of approximately 1 indicates that density around the point is comparable to its neighbors. Scores significantly larger than 1 indicate outliers.

**Value**

A numeric vector of length `ncol(x)` containing LOF values for all data points.

**Author(s)**

Michael Hahsler

**References**

Breunig, M., Kriegel, H., Ng, R., and Sander, J. (2000). LOF: identifying density-based local outliers. In *ACM Int. Conf. on Management of Data*, pages 93-104.

**See Also**

[kNN](#).

**Examples**

```

set.seed(665544)
n <- 100
x <- cbind(
  x=runif(10, 0, 5) + rnorm(n, sd=0.4),
  y=runif(10, 0, 5) + rnorm(n, sd=0.4)
)

### calculate LOF score
lof <- lof(x, k=4)

### distribution of outlier factors
summary(lof)
hist(lof, breaks=10)

### point size is proportional to LOF
plot(x, pch = ".", main = "LOF (k=4)")
points(x, cex = (lof-1)*3, pch = 1, col="red")
text(x[lof>2,], labels = round(lof, 1)[lof>2], pos = 1)

```

---

optics

---

*OPTICS*


---

**Description**

Implementation of the OPTICS (Ordering points to identify the clustering structure) clustering algorithm using a kd-tree.

**Usage**

```

optics(x, eps, minPts = 5, eps_cl,
  search = "kdtree", bucketSize = 10,
  splitRule = "suggest", approx = 0)

optics_cut(x, eps_cl)

```

**Arguments**

x	a data matrix or a distance matrix.
eps	upper limit of the size of the epsilon neighborhood.
minPts	number of minimum points in the eps region (for core points). Default is 5 points.
eps_cl	Threshold to identify clusters (eps_cl <= eps).
search	nearest neighbor search strategy (one of "kdtree", "linear" search, or precomputed "dist") Using precomputed distances is better for high-dimensional, but smaller data sets.
bucketSize	max size of the kd-tree leafs.

splitRule	rule to split the kd-tree. One of "STD", "MIDPT", "FAIR", "SL_MIDPT", "SL_FAIR" or "SUGGEST" (SL stands for sliding). "SUGGEST" uses ANNs best guess.
approx	relative error bound for approximate nearest neighbor searching.

### Details

This implementation of OPTICS implements the original algorithm as described by Ankers et al (1999). OPTICS is similar to DBSCAN, however, for OPTICS eps is only an upper limit for the neighborhood size used to reduce computational complexity. Similar to DBSCAN, minPts is often set to be dimensionality of the data plus one.

OPTICS linearly orders the data points such that points which are spatially closest become neighbors in the ordering. The closest analog to this ordering is dendrogram in single-link hierarchical clustering. The algorithm also calculates the reachability distance for each point. `plot()` produces a reachability-plot which shows each points reachability distance where the points are sorted by OPTICS. Valleys represent clusters (the deeper the valley, the more dense the cluster) and high points indicate points between clusters.

If `eps_cl` is specified, then an algorithm to extract clusters (see Ankers et al, 1999) is used. That is, it internally calls `optics_cut` to extract the clustering. The resulting clustering is similar to what DBSCAN would produce. The only difference is that OPTICS is not able to assign some border points and reports them instead as noise.

See [kNN](#) for more information on the other parameters related to nearest neighbor search.

### Value

An object of class 'optics' with components:

eps	value of eps parameter.
minPts	value of minPts parameter.
order	optics order for the data points in x.
reachdist	reachability distance for each data point in x.
coredist	core distance for each data point in x.

If `eps_cl` was specified or `optics_cut` was called, then in addition the following components are available:

eps_cl	reachability distance for each point in x.
cluster	assigned cluster labels in the order of the data points in x.

### Author(s)

Michael Hahsler

### References

Mihael Ankerst, Markus M. Breunig, Hans-Peter Kriegel, Joerg Sander (1999). OPTICS: Ordering Points To Identify the Clustering Structure. ACM SIGMOD international conference on Management of data. ACM Press. pp. 49–60.

**See Also**

[dbscan](#) in **fpc**.

**Examples**

```
set.seed(2)
n <- 400

x <- cbind(
  x = runif(4, 0, 1) + rnorm(n, sd=0.1),
  y = runif(4, 0, 1) + rnorm(n, sd=0.1)
)

plot(x, col=rep(1:4, time = 100))

### run OPTICS
res <- optics(x, eps = 1, minPts = 10)
res

### get order
res$order

### plot produces a reachability plot
plot(res)

### identify clusters by cutting the reachability plot (black is noise)
res <- optics_cut(res, eps_cl = .065)
res

plot(res)
plot(x, col = res$cluster+1L)

### re-cutting at a higher eps threshold
res <- optics_cut(res, eps_cl = .1)
res
plot(res)
plot(x, col = res$cluster+1L)

### using OPTICS on a precomputed distance matrix
d <- dist(x)
res <- optics(x, eps = 1, minPts = 10)
plot(res)
```

# Index

## \*Topic **clustering**

- dbscan, [2](#)
- optics, [10](#)

## \*Topic **model**

- dbscan, [2](#)
- frNN, [4](#)
- kNN, [6](#)
- kNNdist, [8](#)
- lof, [9](#)
- optics, [10](#)

## \*Topic **plot**

- kNNdist, [8](#)

DBSCAN (dbscan), [2](#)  
dbscan, [2](#), [3](#), [12](#)

frNN, [4](#), [7](#)  
frnn (frNN), [4](#)

kNN, [3](#), [5](#), [6](#), [8](#), [9](#), [11](#)  
knn (kNN), [6](#)  
kNNdist, [8](#)  
kNNdistplot, [3](#)  
kNNdistplot (kNNdist), [8](#)

LOF (lof), [9](#)  
lof, [9](#)

OPTICS (optics), [10](#)  
optics, [10](#)  
optics\_cut (optics), [10](#)